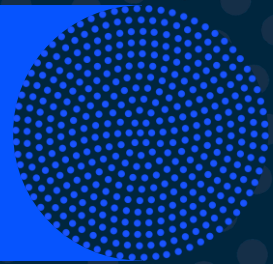


Why use GPUs instead of CPUs?



Short Answer:

They are faster.  
**Way** faster.

The “fastest” and “largest” **CPUs** today have speeds around **4 GHs** and about **128 cores**, and costs about **\$5,000**

An “average” GPU like the RTX 3080, has core speeds of **1.5 GHs** and **5,888 cores**, and costs about **\$1,000**

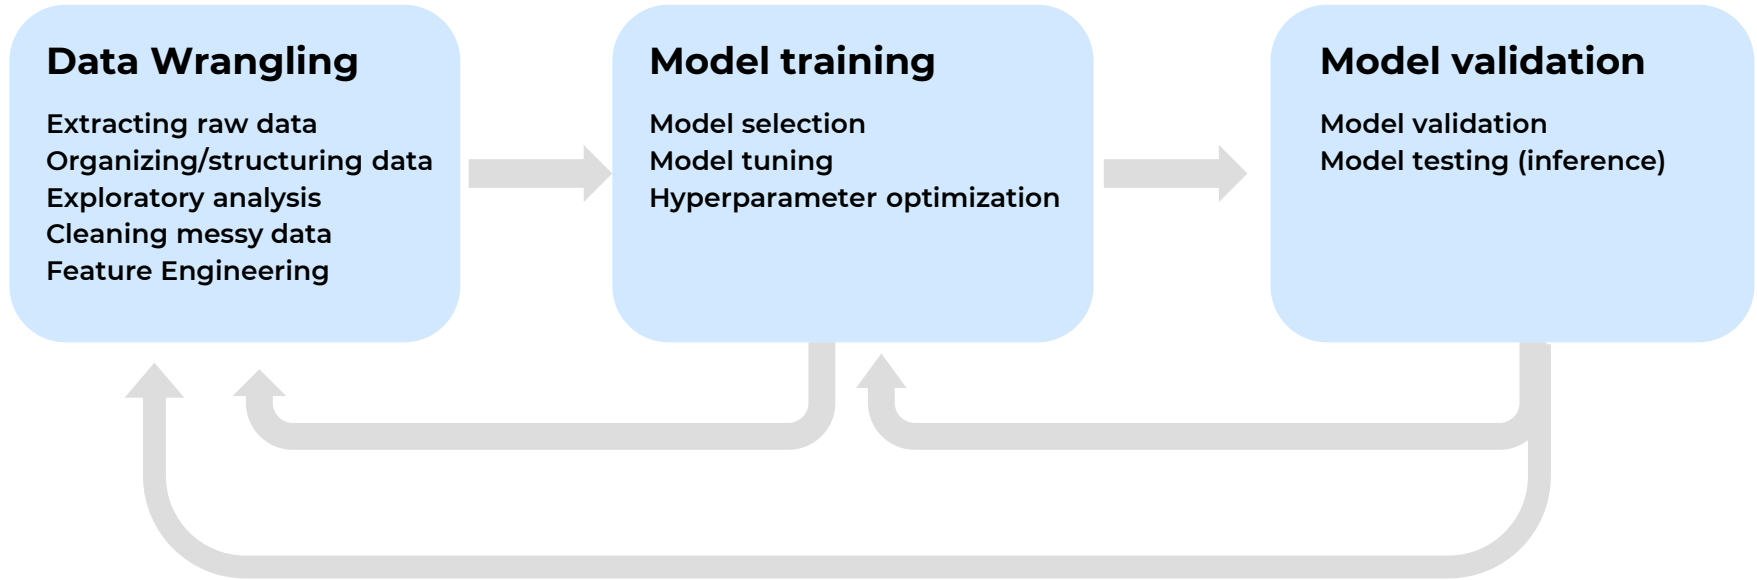
Size **not** to scale

Cycles per second

8,832

512

# Typical Data Science Workflows



# GPU Acceleration

GPU acceleration is commonly concentrated in the modeling portions.

**Data Wrangling**

**Model training**

**Model validation**

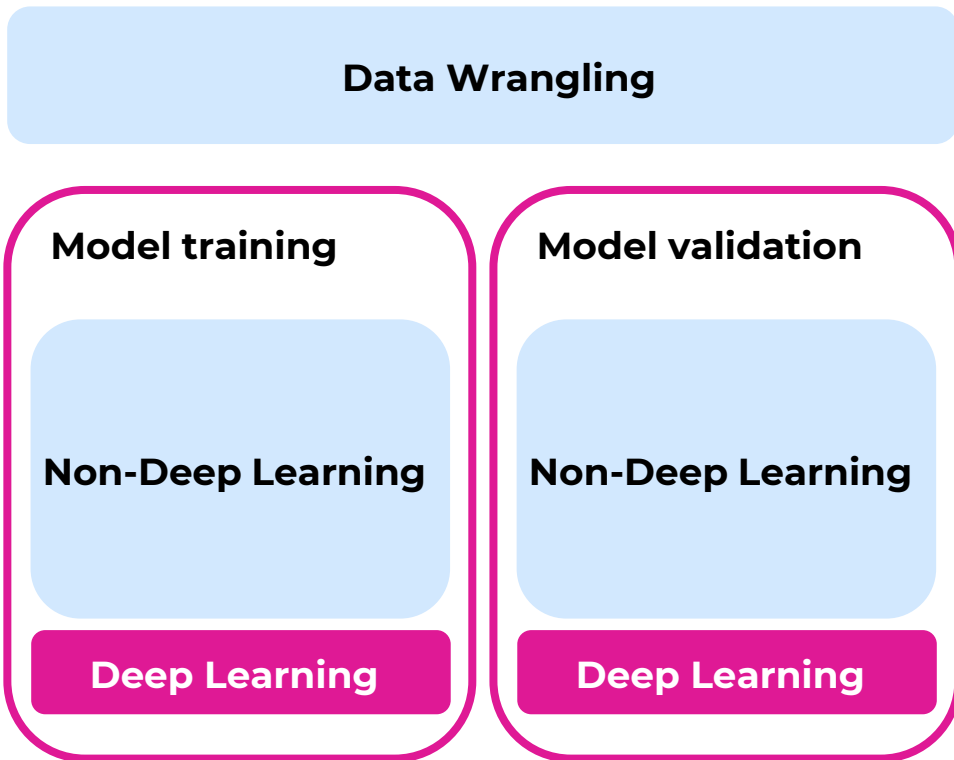
# GPU Acceleration

GPU acceleration is commonly concentrated in the modeling portions.

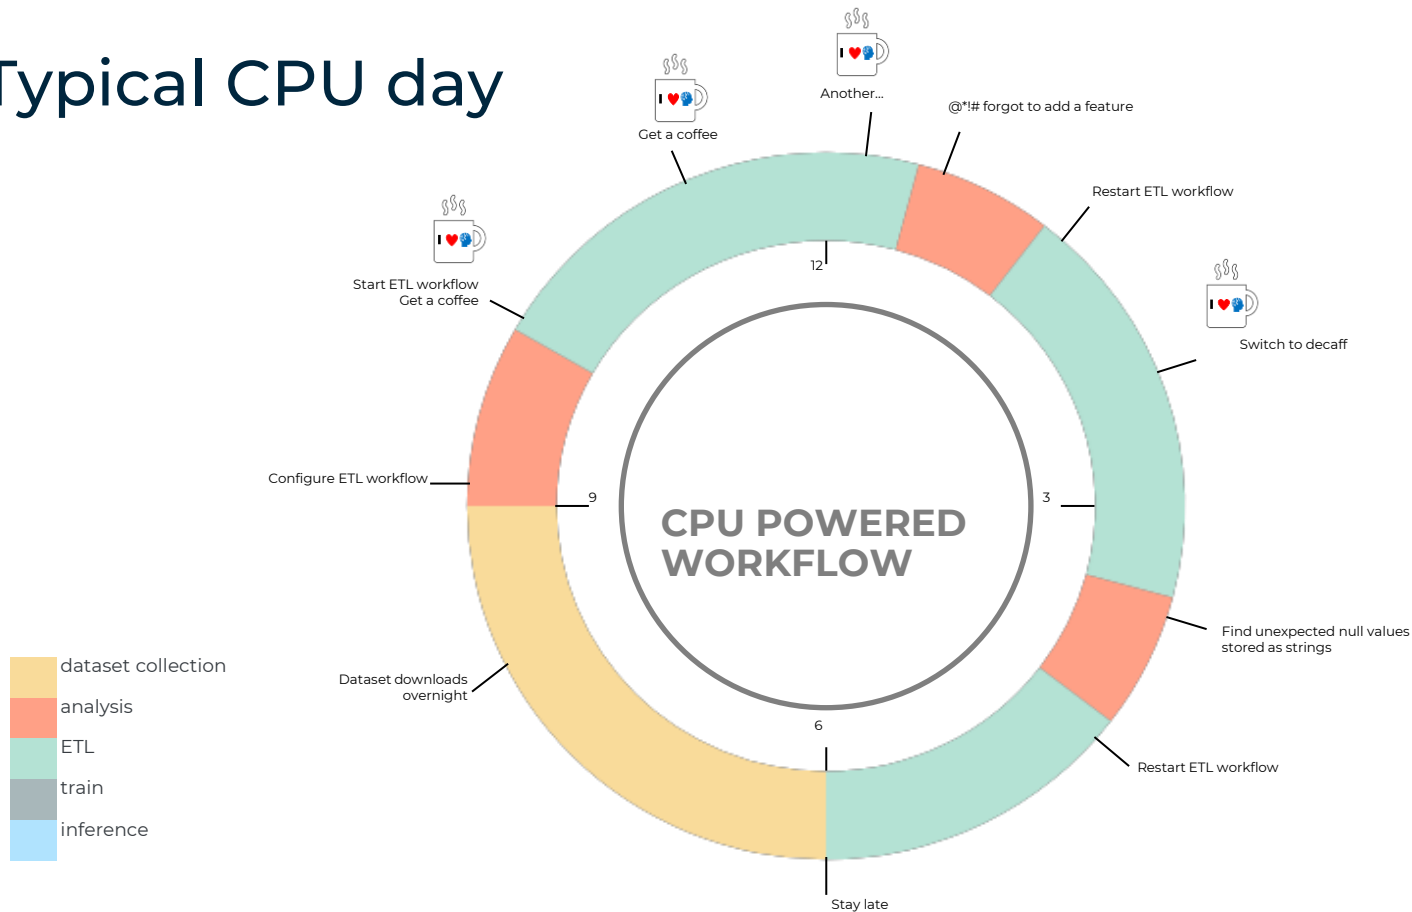
Furthermore, that **concentration** tends to be in the **Deep Learning** and **Neural Networks** subfields.

This isn't necessarily a bad thing; training a model in a matter of hours or days versus weeks is a huge improvement.

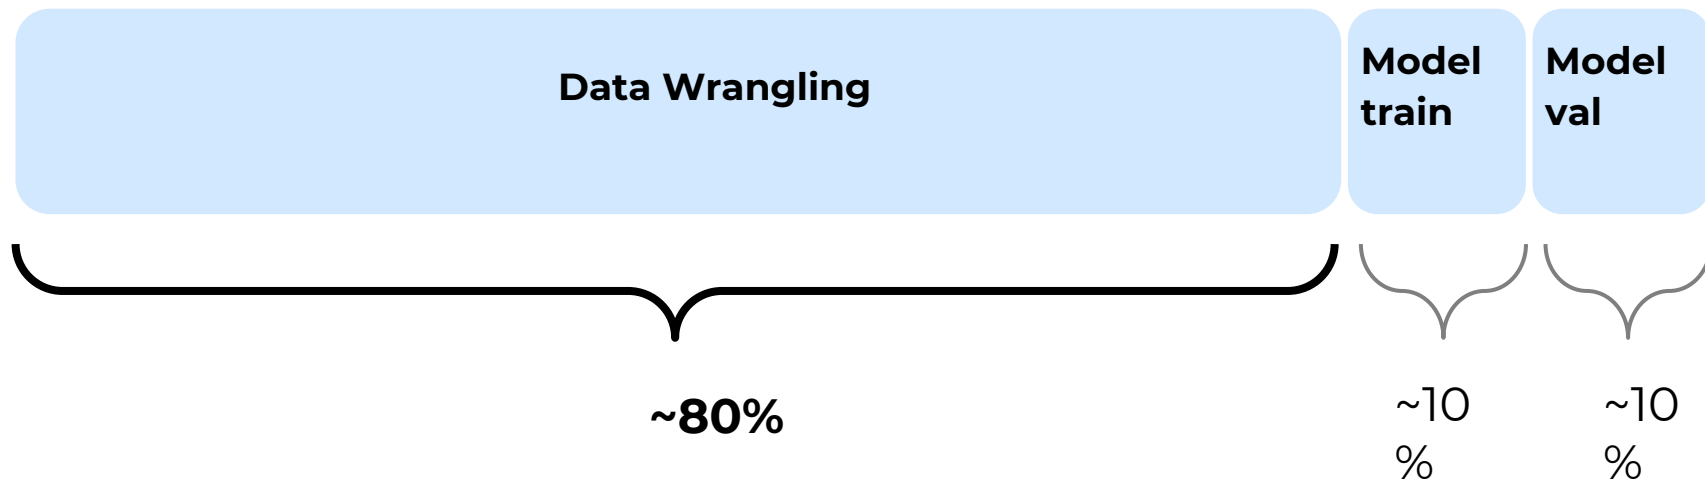
Still, it would be nice to be able to **leverage the GPU for non-deep learning tasks....**



# Typical CPU day

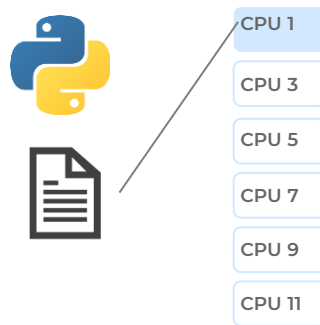
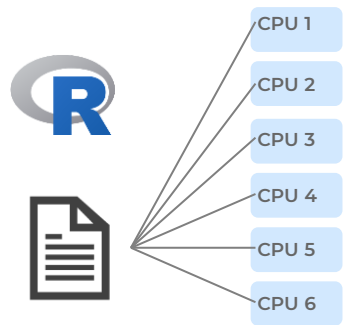


# Points of frustration



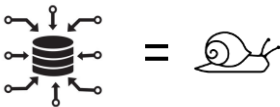
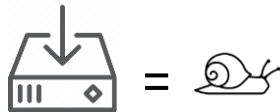
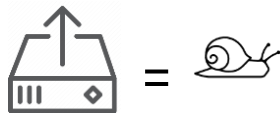
At least **80% of my time** was spend wrestling with the data!  
Very little was left for model training and validation!  
A huge bottleneck in productivity!

# Points of frustration



Some languages, like R, will use all available CPU cores. Python, by purposeful design, is limited to a single CPU.

I can get around this by using python's multiprocessing module, which *mimics* multiprocessing, but it's still an extra step.



Pandas has notoriously **slow** read, write, and aggregation speeds



# A wonderful unified solution

## RAPIDS

The **Rapids.ai** is a suite of GPU accelerated libraries for python and include:

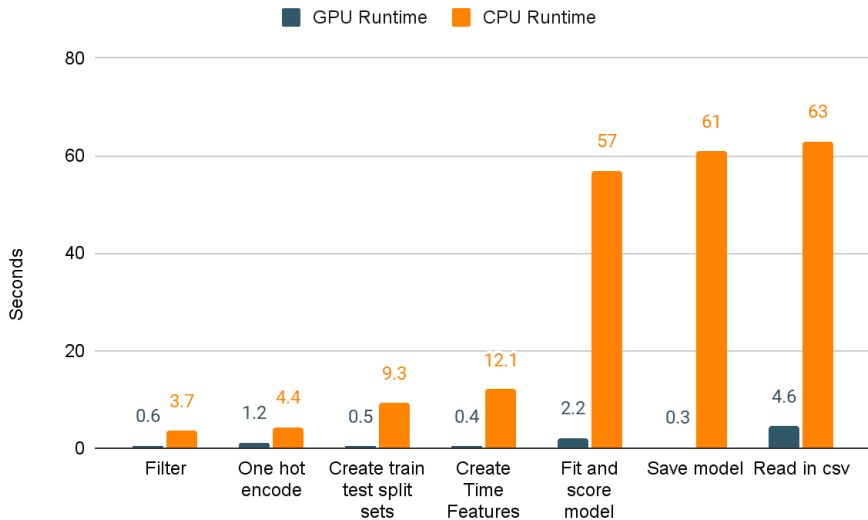
- **cuDF**: mirroring Pandas
- **cuML**: mirroring Scikit-learn
- **cuSignal**: mirroring SciPy
- **cuXFilter**: framework for visualization
- **cuGraph**: mirroring NetworkX
- and more!

In addition **XGBoost**, **Dask**, and **Spark** all have GPU implementations that utilize the Rapids engine.

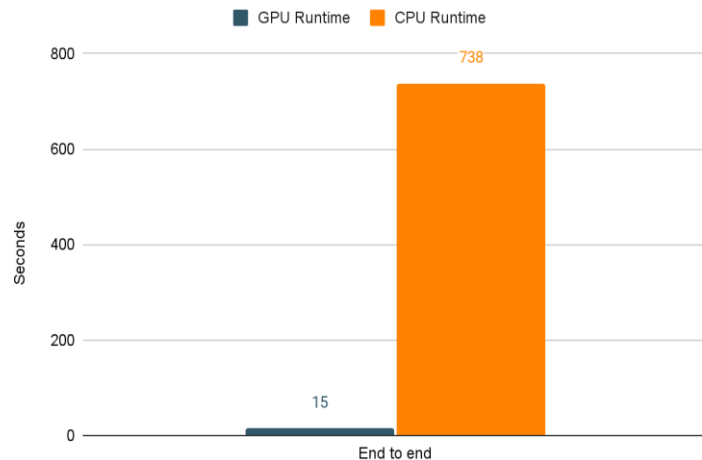
# Comparing speeds: trip data for NYC taxis

I came across these on the internet...

## Individual Tasks



## Entire pipeline



...and had to test it out myself....

# Benchmark: 20M rows & 10 columns

On CPU, in the time that it took to:

Load data

Write date

The GPU accomplished:

Load data

Write date

Describe dataframe

Set index on dataframe

Concat multiple dataframes

Groupby aggregation (mean)

Fit label encoder

Encode data

Scale data

split data

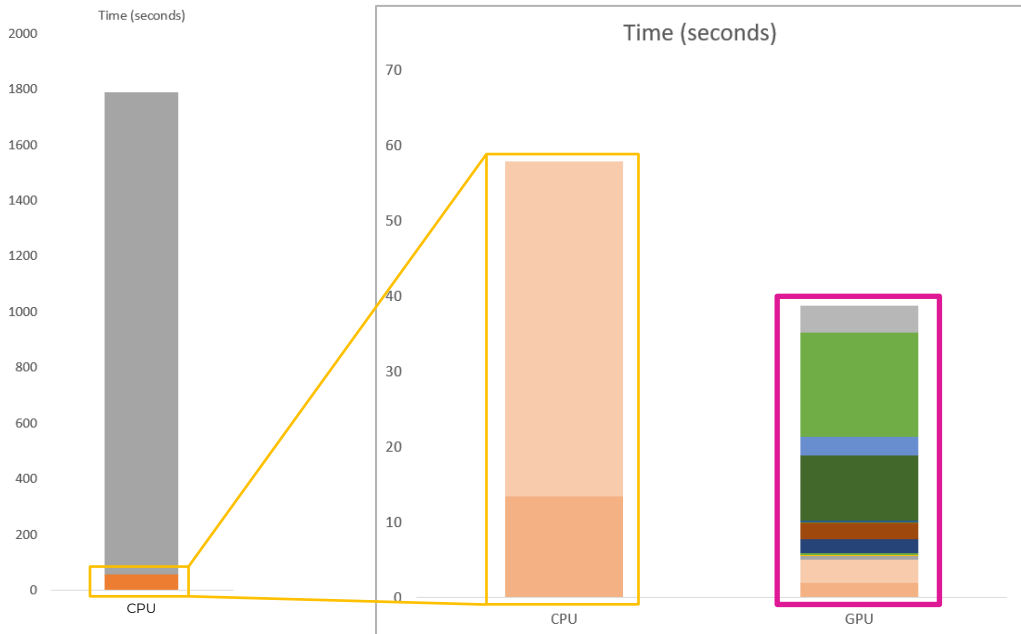
OLS Regression

Logistic Regression

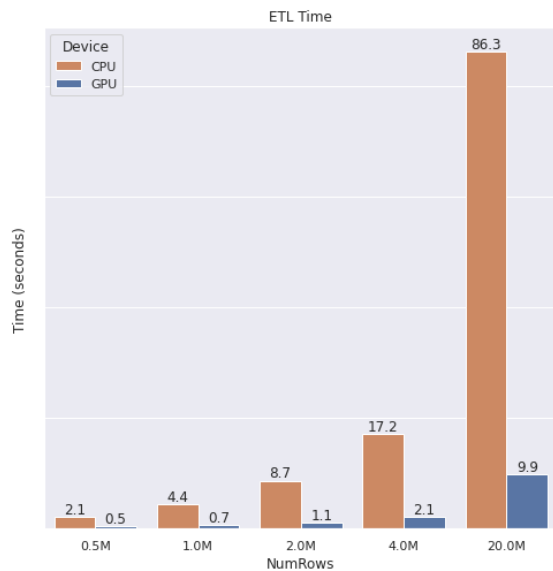
K-Means

Random Forest Classifier

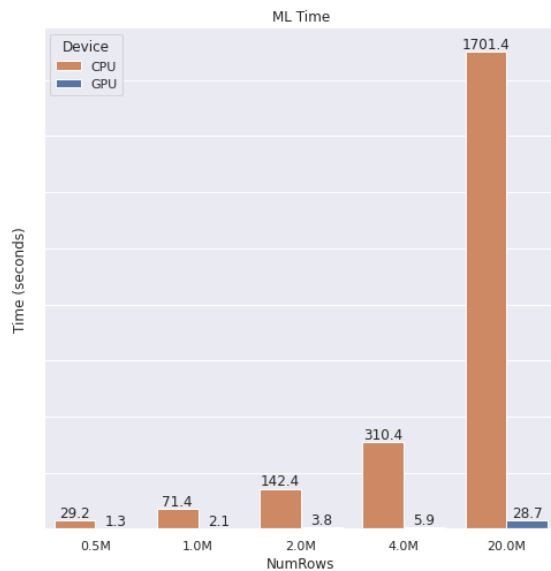
Gradient Boosting



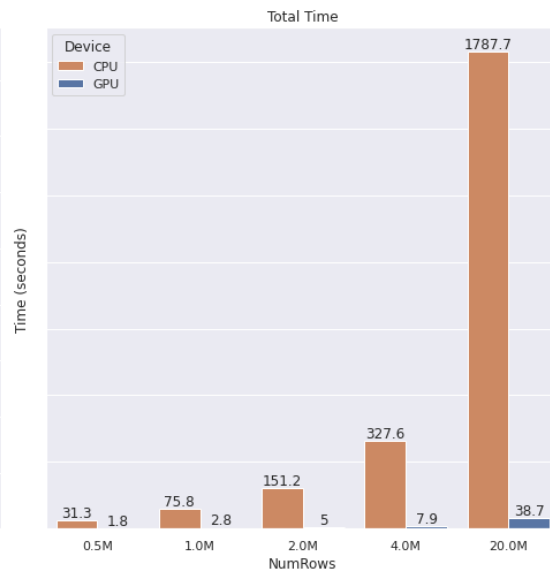
# Benchmark: across different file sizes



For **ETL**, we see more than **85% reduction in time**



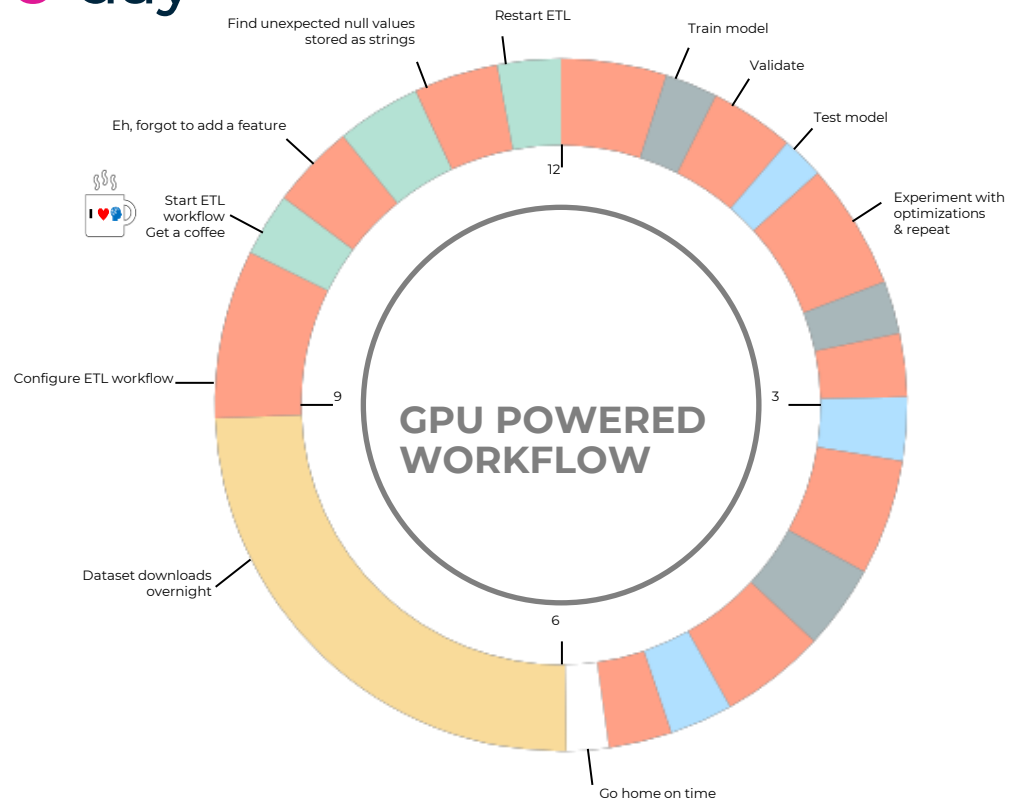
For **ML**, we see a more than **98% reduction in time**



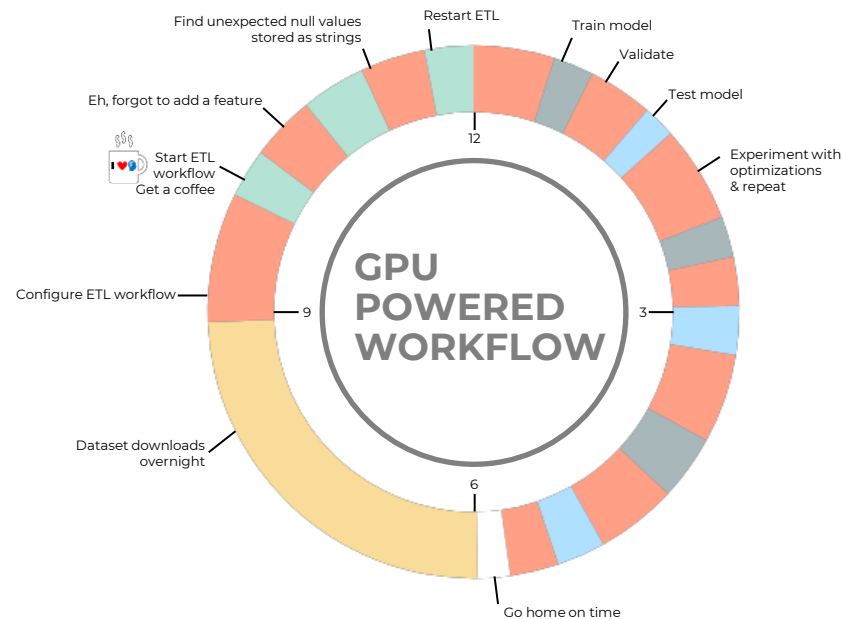
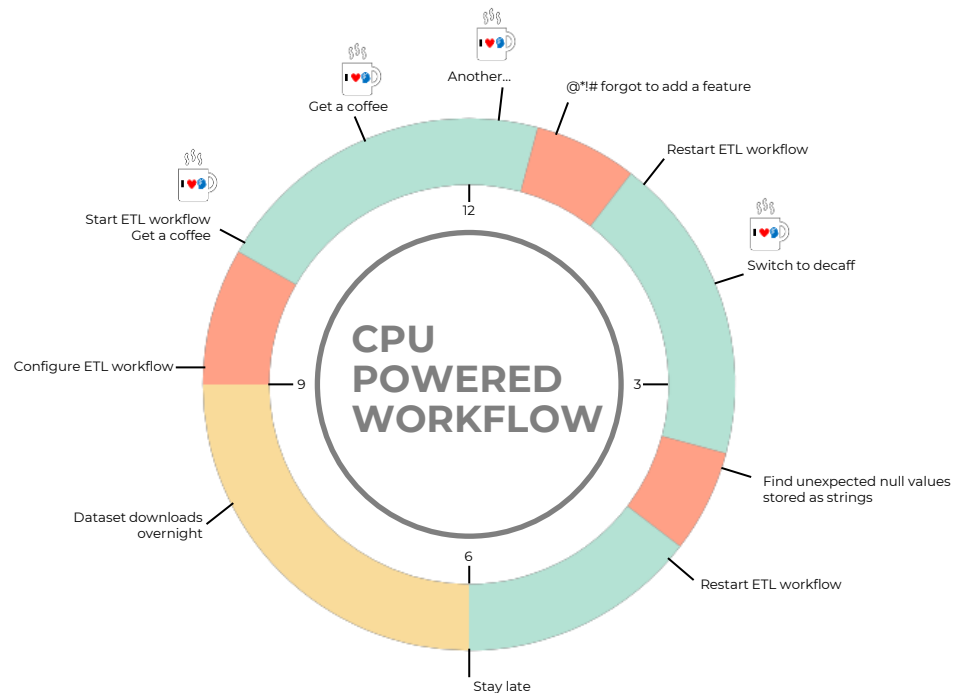
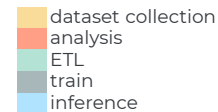
**End-to-end**, we see a more than **95% reduction in time**

The speed gain are **astounding!**

# Typical GPU day

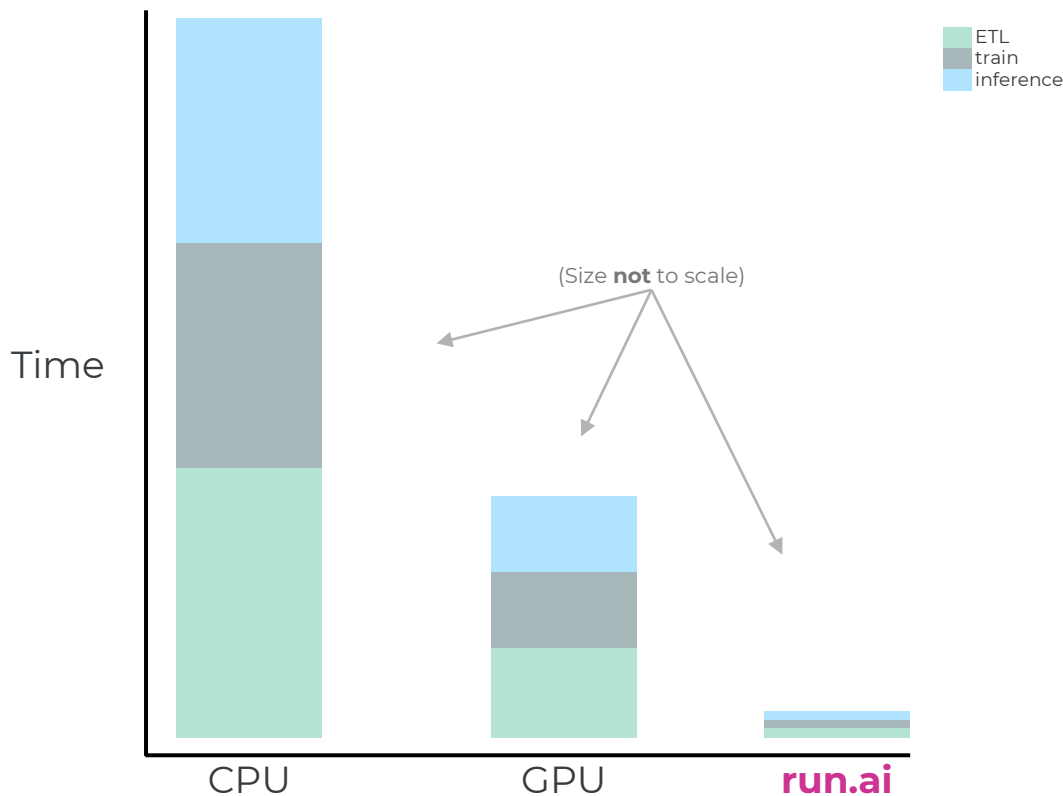


# Comparing CPU vs GPU day

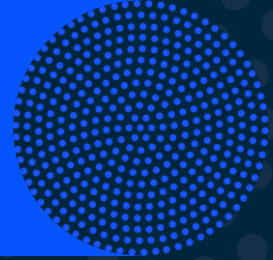


# Speed gains with **run.ai**

Using **run.ai**'s state-of-the-art **scheduler** and **fractional GPU** capabilities, speed gains can be pushed to their maximum



## FAQs





## FAQs

All of my code is already in Pandas and Scikit-learn. Do I have to convert it all?

A valid concern!

cuDF was specifically build to mirror Pandas.

cuML was specifically build to mirror Scikit-Learn.

In nearly all cases, your Pandas/Scikit-Learn code is the same as the cuDF/cuML code!

## FAQs

What if there's something I *must* use pandas for?

No problem!

It's easy to convert back and forth between Pandas Dataframes and cuDF Dataframes!

```
# from cuDF to Pandas  
df = df.to_pandas()
```

```
# your unique function  
df = my_function(df)
```

```
# from Pandas to cuDF  
df = cudf.from_pandas(df)
```

## FAQs

### What if my data is too large for GPU memory?

No problem!

The Rapids ecosystem includes [Dask-GPU](#), which is designed to handle extremely large datasets, in a distributed way.

Best of all, Dask was originally designed to be a distributed extension of Pandas, so implementing it will be very easy if you are used to pandas!

If you are used to using [Spark](#), Apache has also incorporated the Rapids.ai engine to leverage GPU for ETL pipelines.

# FAQs

GPU instances are expensive.  
Wouldn't it be cheaper to run a large CPU cluster rather than a GPU?

Actually, because a GPU can process data so much more quickly, the net cost of running a workflow on a GPU vs a large CPU cluster is actually less!

Rapids Accelerator for Apache Spark reaps the benefit of GPU performance while saving infrastructure costs.



\*ETL for FannieMae Mortgage Dataset (~200GB) as shown in our [demo](#). Costs based on Cloud T4 GPU instance market price & V100 GPU price on Databricks Standard edition

Use GPUs!

